

TTIC 31230, Fundamentals of Deep Learning

David McAllester, Winter 2020

Shannon's Source Coding Theorem

Huffman Coding

Perils of Continuous Information Theory

Source Coding Theorem

Consider a probability distribution Pop on a finite set S .

Consider a code C assigning a bit string code word $C(y_1, \dots, y_B)$ to each possible batch of B elements with $y_i \sim \text{Pop}$.

Source coding theorem: As $B \rightarrow \infty$ the optimal coding uses exactly $H(\text{Pop})$ bits per batch element.

Prefix Free Codes

Let S be a finite set.

Let C be assignment of a bit string $C(y)$ to each $y \in S$.

C is called *prefix-free* if for $x \neq y$ we have that $C(x)$ is not a prefix of $C(y)$.

A concatenation of sequence of prefix-free code words can be uniquely segmented (parsed) back into a sequence of code words.

Prefix-Free Codes as Trees and as Probabilities

A prefix-free code defines a binary branching tree — branch on the first code bit, then the second, and so on.

The leaves of this tree are labeled with the elements of S .

The code defines a probability distribution on S by randomly selecting branches.

We have $P_C(y) = 2^{-|C(y)|}$.

The Source Coding Theorem

(1) There exists a prefix-free code C such that

$$|C(y)| \leq (-\log_2 \text{Pop}(y)) + 1$$

and hence

$$E_{y \sim \text{Pop}} |C(y)| \leq H(\text{Pop}) + 1$$

(2) For any prefix-free code C

$$E_{y \sim \text{Pop}} |C(y)| = E_{y \sim \text{Pop}} -\ln P_C(y) = H(\text{Pop}, P_C) \geq H(\text{Pop})$$

Code Construction

We construct a code by iterating over $y \in S$ in order of decreasing probability (most likely first).

For each y select a code word $C(y)$ (a tree leaf) with length (depth)

$$|C(y)| = \lceil -\log_2 \text{Pop}(y) \rceil$$

and where $C(y)$ is not an extension of (under) any previously selected code word.

Code Existence Proof

At any point before coding all elements of S we have

$$\sum_{y \in \text{Defined}} 2^{-|C(y)|} \leq \sum_{y \in \text{Defined}} \text{Pop}(y) < 1$$

Therefore there exists an infinite descent into the tree that misses all previous code words.

Hence there exists a code word $C(x)$ not under any previous code word with $|C(x)| = \lceil -\log_2 \text{Pop}(y) \rceil$.

Furthermore $C(x)$ is at least as long as all previous code words and hence $C(x)$ is not a prefix of any previously selected code word.

Huffman Coding Produces an Optimal Code For Finite Distributions

Maintain a list of trees T_1, \dots, T_N where each leaf of each tree is labeled with a population element y .

We will write $y \in T_i$ to mean that some leaf of T_i is labeled with y and for $y \in T_i$ write $d(y, T_i)$ for the depth of the leaf labeled with y in the tree T_i .

Initially each tree is just one root node labeled with a single y value and every y value labels some tree.

Tree Weight

We define the weight of a tree T_i by

$$W(T_i) = \sum_{y \in T_i} \text{Pop}(y) d(y, T_i)$$

The Huffman coding algorithm repeatedly merges the two trees of lowest weight into a single tree until all trees are merged.

When all trees are merged the weight of the final tree is the expected code length.

Optimality of Huffman Coding

Theorem: The Huffman code T for Pop gives an optimal code C — for any other tree T' defining code C' we have

$$E_{y \sim \text{Pop}} |C(y)| \leq E_{y \sim \text{Pop}} |C'(y)|$$

Invariant: The merge operation preserves the invariant that there exists an optimal tree including all the subtrees on the list.

The Merge Operation Preserves the Invariant

Assume there exists an optimal tree containing the given subtrees.

Consider the two subtrees T_i and T_j of minimal weight. Without loss of generality we can assume that T_i is at least as deep as T_j .

Swapping the sibling of T_i for T_j brings T_i and T_j together. This can only improve the average depth (next slide).

Why The Swap Can Only Improve the Tree

We can swap a heavier deeper tree for a shallower lighter tree.

For a subtree T_i of an optimal tree T let $d(T_i)$ be the depth of T_i in T .

The contribution of T_i to the average code length is $d(T_i)W(T_i)$.

For $d(T_1) \geq d(T_2)$ and $W(T_1) \geq W(T_2)$ we have

$$d(T_1)W(T_1) + d(T_2)W(T_2) \geq d(T_1)W(T_2) + d(T_2)W(T_1)$$

Perils of Differential Entropy

Consider a continuous density $p(x)$. For example

$$p(x) = \frac{1}{\sqrt{2\pi} \sigma} e^{\frac{-x^2}{2\sigma^2}}$$

Differential entropy is defined as

$$H(p) \doteq \int \left(\ln \frac{1}{p(x)} \right) p(x) dx = E_{x \sim p} - \ln p(x)$$

Perils of Differential Entropy

$$\begin{aligned} H(\mathcal{N}(0, \sigma)) &= + \int \left(\ln(\sqrt{2\pi}\sigma) + \frac{x^2}{2\sigma^2} \right) p(x) dx \\ &= \ln(\sigma) + \ln(\sqrt{2\pi}) + \frac{1}{2} \end{aligned}$$

$$\lim_{\sigma \rightarrow 0} H(N(0, \sigma)) = -\infty$$

.

Hence differential entropy then depends on the choice of units — a distributions on lengths will have a different entropy when measuring in inches than when measuring in feet.

Differential Cross Entropy can Diverge to $-\infty$

Consider the unsupervised training object.

$$\Phi^* = \operatorname{argmin}_{\Phi} E_{y \sim \text{train}} - \ln p_{\Phi}(y)$$

The training set is finite (discrete).

For each y the density $p_{\Phi}(y)$ can go to infinity.

This will drive the cross entropy training loss to $-\infty$.

Differential Cross Entropy can Diverge to $-\infty$

$$\Phi^* = \operatorname{argmin}_{\Phi} E_{y \sim \text{train}} - \ln p_{\Phi}(y)$$

To avoid divergence to $-\infty$ we can enforce an upper bound on the density $p_{\Phi}(y)$.

We will see how this is implicitly done in continuous variational auto-encoders (VAEs).

Differential Entropy is Actually Infinite

An actual real number carries an infinite number of bits.

Consider quantizing the real numbers into bins.

A continuous probability density p assigns a probability $p(B)$ to each bin.

As the bin size decreases toward zero the entropy of the bin distribution increases toward ∞ .

A meaningful convention is that $H(p) = +\infty$ for any continuous density p .

Differential KL-divergence is Meaningful

$$KL(p, q) = \int \left(\ln \frac{p(x)}{q(x)} \right) p(x) dx$$

This integral can be computed by dividing the real numbers into bins and computing the KL divergence between the distributions on bins.

The KL divergence between the bin distribution typically approaches a finite limit as the bin size goes to zero.

Unlike entropy, differential KL divergence is always non-negative. But as in the discrete case, it can be infinite.

Mutual Information

For two random variables x and y there is a distribution on pairs (x, y) determined by the population distribution.

Mutual information is a KL divergence and hence differential mutual information is meaningful.

$$\begin{aligned} I(x, y) &\doteq KL(p(x, y), p(x)p(y)) \\ &= E_{x,y} \ln \frac{p(x, y)}{p(x)p(y)} \end{aligned}$$

The Data Processing Inequality

For continuous y and z with $z = f(y)$ we get that $H(z)$ can be either larger or smaller than $H(y)$ (consider $z = ay$ for $a > 1$ vs. $a < 1$).

However, mutual information is a KL divergence and is more meaningful than entropy and for $z = f(y)$ we do have

$$I(x, z) \leq I(x, y)$$

END